Garvit Audichya

+91-6375183315 | audichyagarvit010@gmail.com | github | linkedin

PROFILE SUMMARY

Aspiring AI Engineer specializing in machine learning, NLP, and generative AI. Experienced in building LLM-powered systems, RAG pipelines, and evaluation-driven AI workflows. Strong foundation in Python, deep learning, model inference, and LLMOps with hands-on production experience. Proven ability to deliver scalable, reliable GenAI solutions through real-world and hackathon-winning projects.

EXPERIENCE

AI DEVELOPER INTERN | Dec 2025-Present

Debales AI - Remote

- Developing **FastAPI** endpoints to expose client-specific AI agent workflows powered by **LangGraph-based** agent pipelines.
- Working with **MongoDB-backed prompt management**, enabling dynamic prompt loading, versioning, and per-client customization without redeployments.
- Supporting a **multi-tenant production AI system**, integrating agent routing, tool selection, and backend services with proper error handling and observability.

TECHNICAL PROJECTS

VidSense AI — YouTube Transcript RAG Assistant (Chrome Extension) (Demo Video)

- **Architected and deployed a production-grade Retrieval-Augmented Generation (RAG) system** as the primary **AI Developer** for real-time conversational Q&A over YouTube videos, integrating transcript ingestion, semantic chunking, dense vector embeddings, FAISS similarity search, and Hugging Face LLMs, improving answer grounding and reducing hallucinations by **40–50%** compared to naive LLM prompting.
- **Developed a memory-aware, multi-turn reasoning pipeline** with LLM-based query rewriting, session-level conversational memory, and conditional web search augmentation, enabling robust handling of follow-up questions and incomplete transcripts while maintaining **sub-second retrieval latency** and **top-K context precision >90%** on evaluated queries.
- **Implemented LLMOps and evaluation workflows using RAGAS**, tracking context relevance, faithfulness, and answer relevancy; achieved **Context Relevance ≈ 1.0** and **Answer Relevancy ≈ 0.75+** on benchmark queries, with structured logging, source attribution, and modular model abstraction enabling rapid experimentation across open-source LLMs.

Creator Catalyst- AI Content Repurposing Engine (Demo Video)

- Built a **multi-modal AI content automation system** using **Python, Streamlit, Gemini (video analysis & NLP), and FFmpeg** to repurpose long-form videos into multiple publish-ready formats, with structured validation to ensure reliable outputs.
- Won **1st Place at the AI Demos Hackathon** by delivering an end-to-end automated workflow that generated **5+ production-ready content assets per video**, reducing **creator workload by ~80%** and cutting **time-to-publish from hours to minutes.**

SKILLS

Languages: Python, C++, SQL, JavaScript, HTML, CSS

AI & ML: Machine Learning, Deep Learning, NLP, Generative AI, RAG, Prompt Engineering, LLM Evaluation, Model Inference, Vector Databases, Embeddings, Google Gemini

Frameworks & Tools: PyTorch, Scikit-learn, Hugging Face, LangChain, LangGraph, RAGAS, FastAPI, Flask, Docker, MongoDB, FAISS, Chroma, Pandas, NumPy, Streamlit, Git, GitHub, n8n

EDUCATION

Vellore Institute of Technology

B.Tech in **Computer Science & Engineering** (Specialization in AI & ML) - **9.25 CGPA** (Current) *2023– 2027*

Higher Secondary School Certificate (HSSC) – **93%** *2021-2023*

Secondary School Certificate (SSC) - **96.4 %** *2019-2021*

CERTIFICATIONS & ACHIEVEMENTS

- Basic RAG with LangChain using Ollama - Udemy (*link*)
- GEN AI using IBM Watsonx by IBM Developer Skill Network – IBM (*link*)
- **Winner (1st Place)** – *AI Demos × AI for Content Creators Hackathon;* selected as **Jury's Pick**.